

5 X_Xception_central

April 3, 2023

1 Date: 9 2022

2 Method: Cross_Inception

3 Data: Pavia

4 Results v.05

```
[ ]: # Libraries
import pandas as pd
import numpy as np
import seaborn as sn
from sklearn.decomposition import PCA
```

```
[ ]: # Read dataset Pavia
from scipy.io import loadmat

def read_HSI():
    X = loadmat('Pavia.mat')['pavia']
    y = loadmat('Pavia_gt.mat')['pavia_gt']
    print(f"X shape: {X.shape}\ny shape: {y.shape}")
    return X, y

X, y = read_HSI()
```

X shape: (1096, 715, 102)

y shape: (1096, 715)

```
[ ]: # PCA
def applyPCA(X, numComponents): # numComponents=64
    newX = np.reshape(X, (-1, X.shape[2]))
    print(newX.shape)
    pca = PCA(n_components=numComponents, whiten=True)
    newX = pca.fit_transform(newX)
    newX = np.reshape(newX, (X.shape[0], X.shape[1], numComponents))
    return newX, pca, pca.explained_variance_ratio_
```

```
[ ]: # channel_wise_shift
def channel_wise_shift(X,numComponents):
    X_copy = np.zeros((X.shape[0] , X.shape[1], X.shape[2]))
    half = int(numComponents/2)
    for i in range(0,half-1):
        X_copy[:, :, i] = X[:, :, (half-i)*2-1]
    for i in range(half,numComponents):
        X_copy[:, :, i] = X[:, :, (i-half)*2]
    X = X_copy
    return X

[ ]: # Split the hyperspectral image into patches of size windowSize-by-windowSize
    ↳pixels
def Patches_Creating(X, y, windowSize, removeZeroLabels = True): #
    ↳windowSize=15, 25
    margin = int((windowSize - 1) / 2)
    zeroPaddedX = padWithZeros(X, margin=margin)
    # split patches
    patchesData = np.zeros((X.shape[0] * X.shape[1], windowSize, windowSize, X.
    ↳shape[2]),dtype="float16")
    patchesLabels = np.zeros((X.shape[0] * X.shape[1]),dtype="float16")
    patchIndex = 0
    for r in range(margin, zeroPaddedX.shape[0] - margin):
        for c in range(margin, zeroPaddedX.shape[1] - margin):
            patch = zeroPaddedX[r - margin:r + margin + 1, c - margin:c +
    ↳margin + 1]
            patchesData[patchIndex, :, :, :] = patch
            patchesLabels[patchIndex] = y[r-margin, c-margin]
            patchIndex = patchIndex + 1
    if removeZeroLabels:
        patchesData = patchesData[patchesLabels>0,:,:,:]
        patchesLabels = patchesLabels[patchesLabels>0]
        patchesLabels -= 1
    return patchesData, patchesLabels
# padding With Zeros
def padWithZeros(X, margin=2):
    newX = np.zeros((X.shape[0] + 2 * margin, X.shape[1] + 2* margin, X.
    ↳shape[2]),dtype="float16")
    x_offset = margin
    y_offset = margin
    newX[x_offset:X.shape[0] + x_offset, y_offset:X.shape[1] + y_offset, :] = X
    return newX

[ ]: # Split Data
from sklearn.model_selection import train_test_split

def splitTrainTestSet(X, y, testRatio, randomState=345):
```

```

X_train, X_test, y_train, y_test = train_test_split(X, y,
↳test_size=testRatio, random_state=randomState,stratify=y)
return X_train, X_test, y_train, y_test

```

```

[ ]: test_ratio = 0.5

# Load and reshape data for training
X0, y0 = read_HSI()
#X=X0
#y=y0

windowSize=5 # accuracy of
# Score for fold 1: loss of 0.34631192684173584; accuracy of 89.49999809265137%

# to test: 7, 9, 13, 15,

width = windowSize
height = windowSize
img_width, img_height, img_num_channels = windowSize, windowSize, 3

input_image_size=windowSize
INPUT_IMG_SIZE=windowSize

dimReduction=3

InputShape=(windowSize, windowSize, dimReduction)

#X, y = loadData(dataset) channel_wise_shift
X1,pca,ratio = applyPCA(X0,numComponents=dimReduction)
X2_shifted = channel_wise_shift(X1,dimReduction) # channel-wise shift
#X2=X1

#print(f"X0 shape: {X0.shape}\ny0 shape: {y0.shape}")
#print(f"X1 shape: {X1.shape}\nX2 shape: {X2.shape}")

X3, y3 = Patches_Creating(X2_shifted, y0, windowSize=windowSize)
Xtrain, Xtest, ytrain, ytest = splitTrainTestSet(X3, y3, test_ratio)

```

```

X shape: (1096, 715, 102)
y shape: (1096, 715)
(783640, 102)

```

```

[ ]: # Compile the model
#incept_model.compile(optimizer='rmsprop', loss='categorical_crossentropy',
↳metrics=['accuracy'])

```

```
[ ]: print()

import warnings
warnings.filterwarnings("ignore")

# load libraries
from keras.initializers import VarianceScaling
from keras.regularizers import l2
from keras.models import Sequential
from keras.layers import Dense
from sklearn import datasets
from sklearn.model_selection import StratifiedKFold
import numpy as np
```

```
[ ]: # 9 classes names

names = ['1. Water', '2. Trees', '3. Asphalt', '4. Self-Blocking Bricks',
         '5. Bitumen', '6. Tiles', '7. Shadows',
         '8. Meadows', '9. Bare Soil']
```

```
[ ]: from tensorflow.keras.applications import EfficientNetB0
from keras.applications import densenet, inception_v3, mobilenet, resnet,
    ↳ vgg16, vgg19, xception
from tensorflow.keras import layers
from keras.layers import Dense, GlobalAveragePooling2D, Dropout, Flatten
import tensorflow as tf

'''
#model = EfficientNetB0(weights='imagenet')

def build_model(num_classes):
    inputs = layers.Input(shape=(windowSize, windowSize, 3))
    #x = img_augmentation(inputs)
    model = xception.Xception(weights='imagenet', include_top=False,
    ↳ input_tensor=inputs)

    #model1 = resnet.ResNet50(weights='imagenet')

    # Freeze the pretrained weights
    model.trainable = False

    # Rebuild top
    x = layers.GlobalAveragePooling2D(name="avg_pool")(model.output)
```

```

x = layers.BatchNormalization()(x)

x = model.output

x = GlobalAveragePooling2D()(x)
# let's add a fully-connected layer
x = Dense(256, activation='relu')(x)
x = Dropout(0.25)(x)

top_dropout_rate = 0.2
#x = layers.Dropout(top_dropout_rate, name="top_dropout")(x)
outputs = layers.Dense(9, activation="softmax", name="pred")(x)

# Compile
model = tf.keras.Model(inputs, outputs, name="EfficientNet")
optimizer = tf.keras.optimizers.Adam(learning_rate=1e-3)
model.compile(
    optimizer=optimizer, loss="categorical_crossentropy",
    metrics=["accuracy"])
return model

```

```

[ ]: '\'\n#model = EfficientNetB0(weights=\'imagenet\')\n\nndef
build_model(num_classes):\n    inputs = layers.Input(shape=(windowSize,
windowSize, 3))\n    #x = img_augmentation(inputs)\n    model =
ception.Xception(weights=\'imagenet\', include_top=False,
input_tensor=inputs)\n\n    #model1 =
resnet.ResNet50(weights=\'imagenet\')\n\n\n    # Freeze the pretrained weights\n
model.trainable = False\n\n    # Rebuild top\n    x =
layers.GlobalAveragePooling2D(name="avg_pool")(model.output)\n    x =
layers.BatchNormalization()(x)\n\n    x = model.output\n\n\n    x =
GlobalAveragePooling2D()(x)\n    # let's add a fully-connected layer\n    x =
Dense(256, activation=\'relu\')(x)\n    x = Dropout(0.25)(x)\n\n\n
top_dropout_rate = 0.2\n    #x = layers.Dropout(top_dropout_rate,
name="top_dropout")(x)\n    outputs = layers.Dense(9, activation="softmax",
name="pred")(x)\n\n    # Compile\n    model = tf.keras.Model(inputs, outputs,
name="EfficientNet")\n    optimizer =
tf.keras.optimizers.Adam(learning_rate=1e-3)\n    model.compile(\n
optimizer=optimizer, loss="categorical_crossentropy", metrics=["accuracy"]\n
)\n    return model\n'

```

```

[ ]: from tensorflow.keras.applications import EfficientNetB0

```

```

def build_model(num_classes):
    inputs = layers.Input(shape=(windowSize, windowSize, 3))
    #x = img_augmentation(inputs)
    #model = EfficientNetB0(include_top=False, input_tensor=inputs,
    ↪weights="imagenet")
    model = xception.Xception(weights='imagenet', include_top=False,
    ↪input_tensor=inputs)

    # Freeze the pretrained weights
    #model.trainable = False

    # Rebuild top
    x = layers.GlobalAveragePooling2D(name="avg_pool")(model.output)
    x = layers.BatchNormalization()(x)

    top_dropout_rate = 0.2
    x = layers.Dropout(top_dropout_rate, name="top_dropout")(x)
    outputs = layers.Dense(9, activation="softmax", name="pred")(x)

    # Compile
    model = tf.keras.Model(inputs, outputs, name="EfficientNet")
    optimizer = tf.keras.optimizers.Adam(learning_rate=1e-3)
    model.compile(
        optimizer=optimizer, loss="categorical_crossentropy",
    ↪metrics=["accuracy"])
    return model

```

```
[ ]: model = build_model(num_classes=9)
```

```

[ ]: def unfreeze_model(model):
    # We unfreeze the top 20 layers while leaving BatchNorm layers frozen
    for layer in model.layers[-20:]:
        if not isinstance(layer, layers.BatchNormalization):
            layer.trainable = True

    optimizer = tf.keras.optimizers.Adam(learning_rate=1e-4)
    model.compile(
        optimizer=optimizer, loss="categorical_crossentropy",
    ↪metrics=["accuracy"])

```

```
[ ]: import matplotlib.pyplot as plt
```

```
def plot_hist(hist):
```

```

plt.plot(hist.history["accuracy"])
plt.plot(hist.history["val_accuracy"])
plt.title("model accuracy")
plt.ylabel("accuracy")
plt.xlabel("epoch")
plt.legend(["train", "validation"], loc="upper left")
plt.show()

```

```

[ ]: from tensorflow.keras.losses import sparse_categorical_crossentropy
from tensorflow.keras.optimizers import Adam
from sklearn.model_selection import KFold
from tensorflow.keras import layers

import numpy as np
from sklearn.metrics import confusion_matrix, accuracy_score, \
    ↪classification_report, cohen_kappa_score
import matplotlib.pyplot as plt
from keras.applications.inception_resnet_v2 import InceptionResNetV2, \
    ↪preprocess_input
from keras.layers import Dense, GlobalAveragePooling2D, Dropout, Flatten
from keras.models import Model

import tensorflow as tf

# configuration
confmat = 0
batch_size = 50
loss_function = sparse_categorical_crossentropy
no_classes = 9
no_epochs = 10
optimizer = Adam()
verbosity = 1
num_folds = 5

NN=len(Xtrain)
NN=700
#NN=5000

input_train=Xtrain[0:NN]
target_train=ytrain[0:NN]

input_test=Xtest[0:NN]
target_test=ytest[0:NN]

# Determine shape of the data
input_shape = (img_width, img_height, img_num_channels)

```

```

# Parse numbers as floats
_train = _train.astype('float32')
_test = _test.astype('float32')

# Normalize data
_train = _train / 255
_test = _test / 255

# Define per-fold score containers
acc_per_fold = []
loss_per_fold = []

Y_pred=[]
y_pred=[]
# Merge inputs and targets
inputs = np.concatenate((_train, _test), axis=0)
targets = np.concatenate((target_train, target_test), axis=0)

# Define the K-fold Cross Validator
kfold = KFold(n_splits=num_folds, shuffle=True)

# K-fold Cross Validation model evaluation
fold_no = 1
for train, test in kfold.split(inputs, targets):

    # model architecture

    # Compile the model
    #model.compile(optimizer='rmsprop', loss='categorical_crossentropy',
    →metrics=['accuracy'])

    # Compile the model
    # model.compile(optimizer='rmsprop', loss='categorical_crossentropy',
    →metrics=['accuracy'])

    model = build_model(num_classes=9)
    #model.compile(loss=loss_function, optimizer='rmsprop',metrics=['accuracy'])

    #model.summary()

    #unfreeze_model(model)
    model.compile(loss=loss_function, optimizer='rmsprop',metrics=['accuracy'])

```



```

# Generate a print
↳
↳ print('-----')
print(f'Training for fold {fold_no} ...')

# Fit data to model
#model.summary()

history = model.fit(inputs[train], targets[train],
                    validation_data = (inputs[test], targets[test]),
                    epochs=no_epochs, verbose=2 )
plt.figure()
plot_hist(history)
# hist = model.fit(inputs[train], targets[train],
#                  steps_per_epoch=(29943/batch_size),
#                  epochs=5,
#                  validation_data=(inputs[test], targets[test]),
#                  validation_steps=(8000/batch_size),
#                  initial_epoch=20,
#                  verbose=1 )
plt.figure()

# Generate generalization metrics
scores = model.evaluate(inputs[test], targets[test], verbose=0)
print(f'Score for fold {fold_no}: {model.metrics_names[0]} of {scores[0]};
↳ {model.metrics_names[1]} of {scores[1]*100}%')
acc_per_fold.append(scores[1] * 100)
loss_per_fold.append(scores[0])

# confusion_matrix
Y_pred = model.predict(inputs[test])
y_pred = np.argmax(Y_pred, axis=1)
#target_test=targets[test]

confusion = confusion_matrix(targets[test], y_pred)
df_cm = pd.DataFrame(confusion, columns=np.unique(names), index = np.
↳ unique(names))
df_cm.index.name = 'Actual'
df_cm.columns.name = 'Predicted'
plt.figure(figsize = (9,9))
sn.set(font_scale=1.4)#for label size
sn.heatmap(df_cm, cmap="Reds", annot=True, annot_kws={"size": 16}, fmt='d')
plt.savefig('cmap.png', dpi=300)
print(confusion_matrix(targets[test], y_pred))

```

```

confmat      = confmat + confusion;

# Increase fold number
fold_no = fold_no + 1

# == average scores ==
print('-----')
print('Score per fold')
for i in range(0, len(acc_per_fold)):
    ↪print('-----')
    ↪print(f'> Fold {i+1} - Loss: {loss_per_fold[i]} - Accuracy:↪
    ↪{acc_per_fold[i]}%')
print('-----')
print('Average scores for all folds:')
print(f'> Accuracy: {np.mean(acc_per_fold)} (+- {np.std(acc_per_fold)})')
print(f'> Loss: {np.mean(loss_per_fold)}')
print('-----')

Overall_Conf = pd.DataFrame(confmat, columns=np.unique(names), index = np.
    ↪unique(names))
Overall_Conf.index.name = 'Actual Overall'
Overall_Conf.columns.name = 'Predicted Overall'
plt.figure(figsize = (10,8))
sn.set(font_scale=1.4)#for label size
sn.heatmap(Overall_Conf, cmap="Reds", annot=True,annot_kws={"size": 16},↪
    ↪fmt='d')
plt.savefig('cmap.png', dpi=300)
print(Overall_Conf)

# Notes for next trial

# window size=25 __> will work
# window size=5 --> Only Bayesian will work
# Need to test (7, 9, 11, 13, 15) window sizes
# When the accuracy is decreasing, it's not right.
# When need to get acc over 0.7

```

Training for fold 1 ...

Epoch 1/10

35/35 - 37s - loss: 1.6848 - accuracy: 0.4848 - val_loss: 1.5543 - val_accuracy:
0.2821 - 37s/epoch - 1s/step

Epoch 2/10

35/35 - 31s - loss: 0.9981 - accuracy: 0.7000 - val_loss: 1.4624 - val_accuracy:

0.5214 - 31s/epoch - 881ms/step

Epoch 3/10

35/35 - 31s - loss: 0.6657 - accuracy: 0.7955 - val_loss: 1.4448 - val_accuracy:

0.5214 - 31s/epoch - 893ms/step

Epoch 4/10

35/35 - 30s - loss: 0.9612 - accuracy: 0.7571 - val_loss: 1.3767 - val_accuracy:

0.5643 - 30s/epoch - 857ms/step

Epoch 5/10

35/35 - 29s - loss: 0.6493 - accuracy: 0.7991 - val_loss: 1.2371 - val_accuracy:

0.6964 - 29s/epoch - 832ms/step

Epoch 6/10

35/35 - 29s - loss: 0.5050 - accuracy: 0.8313 - val_loss: 1.0850 - val_accuracy:

0.7679 - 29s/epoch - 821ms/step

Epoch 7/10

35/35 - 29s - loss: 0.5526 - accuracy: 0.8241 - val_loss: 0.5527 - val_accuracy:

0.8464 - 29s/epoch - 820ms/step

Epoch 8/10

35/35 - 29s - loss: 0.4624 - accuracy: 0.8518 - val_loss: 0.4942 - val_accuracy:

0.8643 - 29s/epoch - 824ms/step

Epoch 9/10


35/35 - 28s - loss: 0.5121 - accuracy: 0.8536 - val_loss: 5.2260 - val_accuracy:

0.8357 - 28s/epoch - 812ms/step

Epoch 10/10

35/35 - 29s - loss: 0.4555 - accuracy: 0.8714 - val_loss: 1.7005 - val_accuracy:

0.8964 - 29s/epoch - 815ms/step



Score for fold 1: loss of 1.7004808187484741; accuracy of 89.64285850524902%

9/9 [=====] - 1s 34ms/step

```
[[134  0  0  0  0  0  0  0  0]
 [  0  7  0  0  3  0  0  0  0]
 [  0  2  0  0  4  1  0  1  0]
 [  0  1  0  0  3  0  0  0  0]
 [  0  3  1  0  8  0  1  1  0]
 [  1  0  0  0  1 15  0  0  0]
 [  0  1  0  0  3  0  5  0  0]
 [  0  1  0  0  1  0  0 77  0]
 [  0  0  0  0  0  0  0  0  5]]
```

Training for fold 2 ...

Epoch 1/10

35/35 - 36s - loss: 1.6127 - accuracy: 0.5000 - val_loss: 1.8610 - val_accuracy:
0.2786 - 36s/epoch - 1s/step

Epoch 2/10

35/35 - 31s - loss: 1.3385 - accuracy: 0.6277 - val_loss: 1.5071 - val_accuracy:
0.5464 - 31s/epoch - 895ms/step

Epoch 3/10

35/35 - 31s - loss: 0.9322 - accuracy: 0.7268 - val_loss: 1.3973 - val_accuracy:
0.5893 - 31s/epoch - 876ms/step

Epoch 4/10

35/35 - 29s - loss: 0.7721 - accuracy: 0.7589 - val_loss: 1.4485 - val_accuracy:
0.6536 - 29s/epoch - 823ms/step

Epoch 5/10

35/35 - 29s - loss: 0.6196 - accuracy: 0.8179 - val_loss: 1.4363 - val_accuracy:
0.6500 - 29s/epoch - 837ms/step

Epoch 6/10

35/35 - 29s - loss: 0.5129 - accuracy: 0.8482 - val_loss: 0.9883 - val_accuracy:
0.7250 - 29s/epoch - 837ms/step

Epoch 7/10

35/35 - 30s - loss: 0.5164 - accuracy: 0.8357 - val_loss: 0.8147 - val_accuracy:
0.7321 - 30s/epoch - 848ms/step

Epoch 8/10

35/35 - 30s - loss: 0.3752 - accuracy: 0.8750 - val_loss: 0.5641 - val_accuracy:
0.8036 - 30s/epoch - 852ms/step

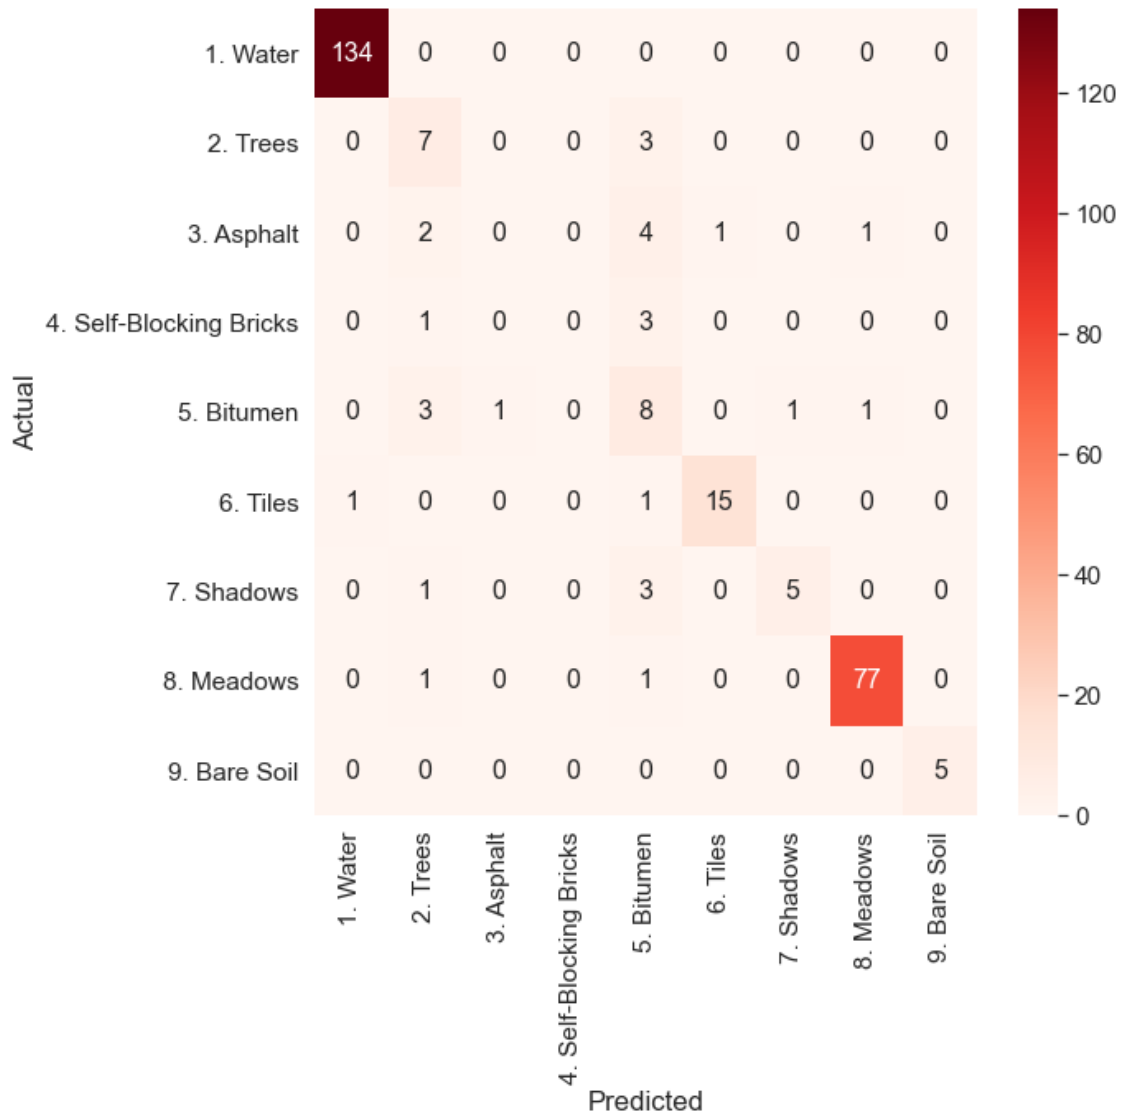
Epoch 9/10


35/35 - 35s - loss: 0.3787 - accuracy: 0.8741 - val_loss: 0.5219 - val_accuracy:
0.7929 - 35s/epoch - 986ms/step

Epoch 10/10

35/35 - 35s - loss: 0.3766 - accuracy: 0.8705 - val_loss: 0.5310 - val_accuracy:
0.8000 - 35s/epoch - 1s/step

<Figure size 432x288 with 0 Axes>





Score for fold 2: loss of 0.5309717059135437; accuracy of 80.0000011920929%
 9/9 [=====] - 1s 40ms/step

```
[[128  0  0  0  0  0  0  0  0]
 [  6  9  2  0  0  0  0  0  0]
 [  3  6  5  0  0  0  0  0  0]
 [  0  4  5  0  0  0  0  0  0]
 [  0  0  3  0  6  0  0  0  0]
 [ 14  1  0  0  0  0  0  0  0]
 [  0  5  1  0  0  0  3  0  0]
 [  5  0  0  0  1  0  0 72  0]
 [  0  0  0  0  0  0  0  0 1]]
```

 Training for fold 3 ...

Epoch 1/10

35/35 - 41s - loss: 1.7158 - accuracy: 0.4223 - val_loss: 1.4761 - val_accuracy:
 0.4964 - 41s/epoch - 1s/step

Epoch 2/10

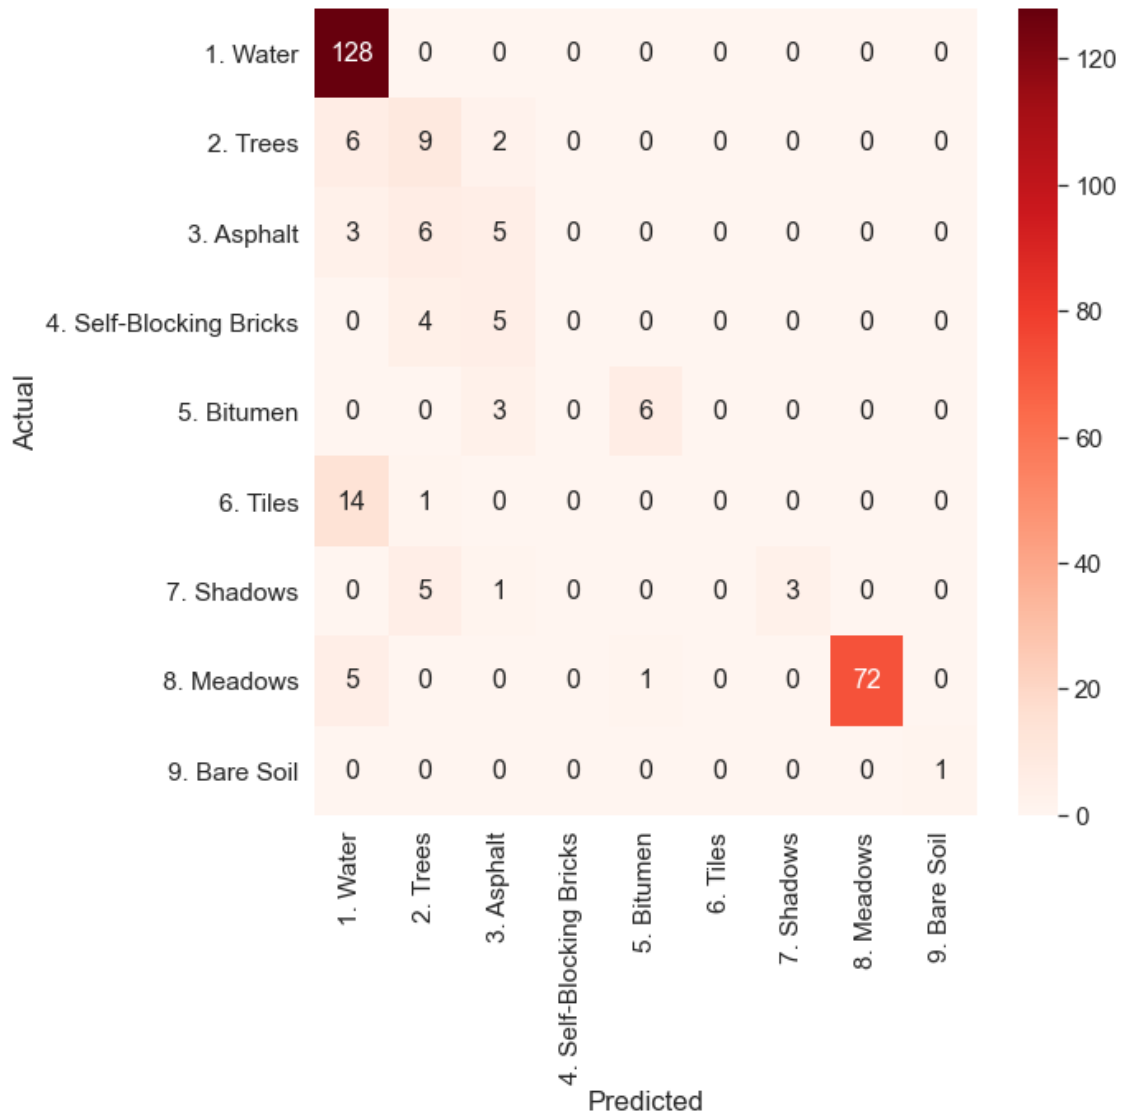
35/35 - 32s - loss: 1.5796 - accuracy: 0.4277 - val_loss: 1.4891 - val_accuracy:
 0.4964 - 32s/epoch - 916ms/step


Epoch 3/10

35/35 - 33s - loss: 1.5800 - accuracy: 0.4277 - val_loss: 1.4784 - val_accuracy:
 0.4964 - 33s/epoch - 940ms/step

Epoch 4/10

35/35 - 32s - loss: 1.5194 - accuracy: 0.4991 - val_loss: 1.4372 - val_accuracy:
0.5321 - 32s/epoch - 921ms/step
Epoch 5/10
35/35 - 34s - loss: 1.0772 - accuracy: 0.6946 - val_loss: 1.4788 - val_accuracy:
0.5107 - 34s/epoch - 958ms/step
Epoch 6/10
35/35 - 35s - loss: 0.8277 - accuracy: 0.7241 - val_loss: 1.4877 - val_accuracy:
0.5750 - 35s/epoch - 1s/step
Epoch 7/10
35/35 - 31s - loss: 0.6614 - accuracy: 0.7946 - val_loss: 2.9608 - val_accuracy:
0.5250 - 31s/epoch - 895ms/step
Epoch 8/10
35/35 - 30s - loss: 0.5993 - accuracy: 0.7893 - val_loss: 5.7760 - val_accuracy:
0.6750 - 30s/epoch - 846ms/step
Epoch 9/10
35/35 - 27s - loss: 0.5942 - accuracy: 0.8107 - val_loss: 1.5897 - val_accuracy:
0.5857 - 27s/epoch - 774ms/step
Epoch 10/10
35/35 - 29s - loss: 0.6350 - accuracy: 0.7795 - val_loss: 1.0341 - val_accuracy:
0.6286 - 29s/epoch - 836ms/step
<Figure size 432x288 with 0 Axes>





Score for fold 3: loss of 1.034110188484192; accuracy of 62.85714507102966%
 9/9 [=====] - 1s 30ms/step

```
[[139  0  0  0  0  0  0  0  0]
 [ 11  0  0  0  2  0  1  0  0]
 [  3  0  0  0  1  0  1  0  0]
 [  1  0  0  0  0  0  1  0  0]
 [  3  0  0  0  2  0  0  2  0]
 [ 24  0  0  0  0  0  0  0  0]
 [ 10  0  0  0  0  0  0  0  0]
 [ 42  0  0  0  0  0  1 31  0]
 [  1  0  0  0  0  0  0  0  4]]
```

 Training for fold 4 ...

Epoch 1/10

35/35 - 32s - loss: 1.5064 - accuracy: 0.6000 - val_loss: 2.4409 - val_accuracy:
 0.3143 - 32s/epoch - 919ms/step

Epoch 2/10

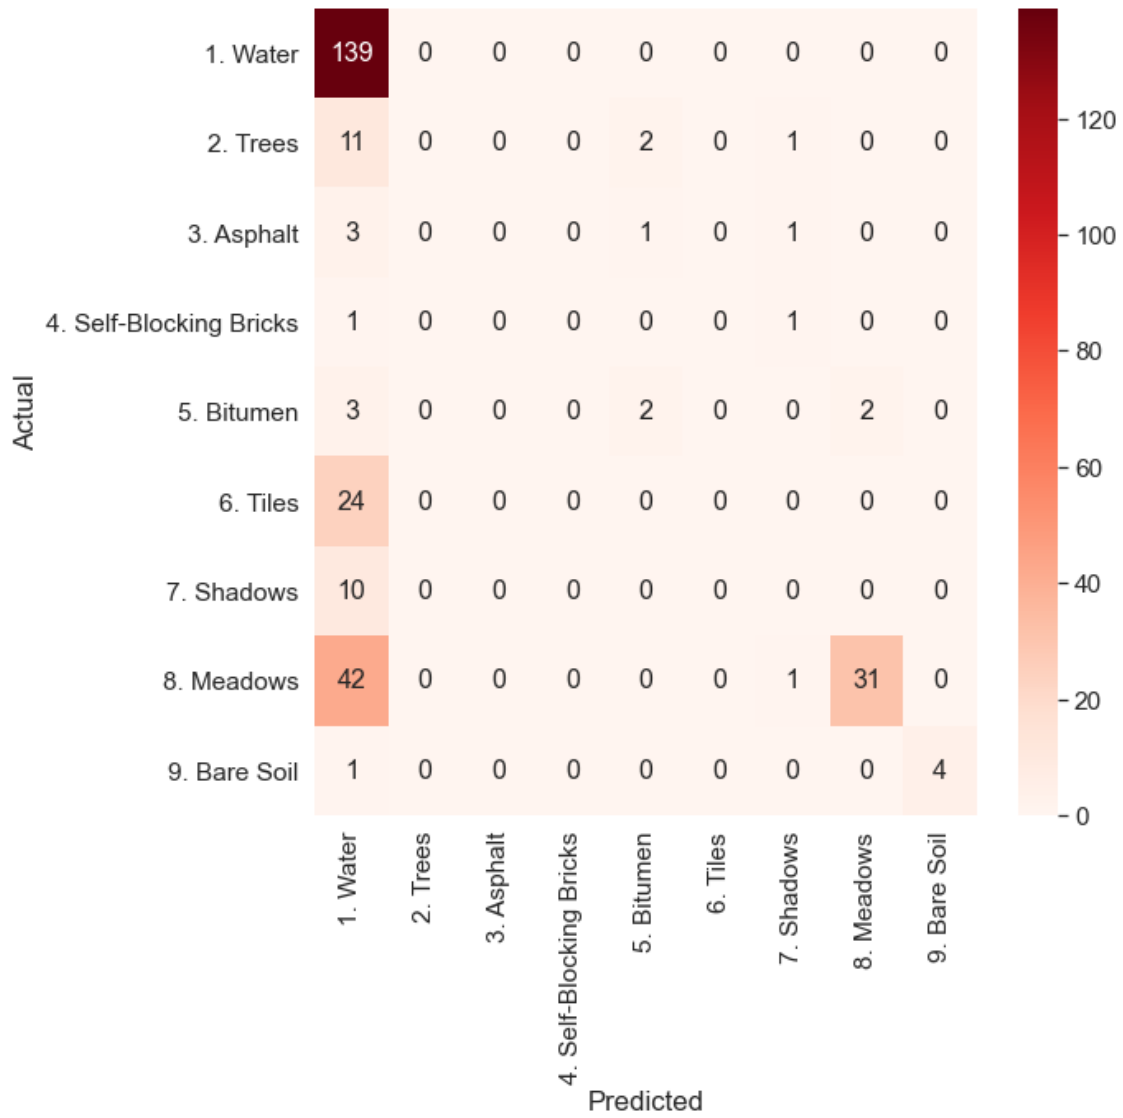
35/35 - 27s - loss: 0.8615 - accuracy: 0.7411 - val_loss: 1.4282 - val_accuracy:
 0.6250 - 27s/epoch - 766ms/step


Epoch 3/10

35/35 - 26s - loss: 0.6919 - accuracy: 0.7946 - val_loss: 1.5018 - val_accuracy:
 0.5357 - 26s/epoch - 744ms/step

Epoch 4/10

35/35 - 27s - loss: 0.6400 - accuracy: 0.8134 - val_loss: 1.1673 - val_accuracy:
0.6714 - 27s/epoch - 765ms/step
Epoch 5/10
35/35 - 24s - loss: 0.5821 - accuracy: 0.8313 - val_loss: 0.9015 - val_accuracy:
0.7286 - 24s/epoch - 699ms/step
Epoch 6/10
35/35 - 24s - loss: 0.5134 - accuracy: 0.8420 - val_loss: 0.7347 - val_accuracy:
0.7857 - 24s/epoch - 688ms/step
Epoch 7/10
35/35 - 25s - loss: 0.4471 - accuracy: 0.8589 - val_loss: 0.5822 - val_accuracy:
0.8214 - 25s/epoch - 710ms/step
Epoch 8/10
35/35 - 27s - loss: 0.4877 - accuracy: 0.8580 - val_loss: 0.4603 - val_accuracy:
0.8643 - 27s/epoch - 759ms/step
Epoch 9/10
35/35 - 25s - loss: 0.4765 - accuracy: 0.8491 - val_loss: 1.7734 - val_accuracy:
0.4893 - 25s/epoch - 702ms/step
Epoch 10/10
35/35 - 24s - loss: 0.4242 - accuracy: 0.8679 - val_loss: 6.3098 - val_accuracy:
0.5071 - 24s/epoch - 698ms/step
<Figure size 432x288 with 0 Axes>





Score for fold 4: loss of 6.309820652008057; accuracy of 50.71428418159485%
 9/9 [=====] - 1s 31ms/step

```
[[11 0 0 0 0 0 0 0 98]
 [ 0 6 1 0 0 4 0 0 0]
 [ 0 2 8 0 0 0 0 0 0]
 [ 0 2 0 4 0 0 0 0 0]
 [ 0 0 2 13 3 0 0 0 0]
 [ 3 0 0 0 0 15 0 0 0]
 [ 0 0 0 0 0 2 6 0 0]
 [ 1 0 0 0 7 0 0 82 0]
 [ 0 0 0 0 0 0 0 3 7]]
```

 Training for fold 5 ...

Epoch 1/10

35/35 - 31s - loss: 1.6524 - accuracy: 0.4455 - val_loss: 1.5890 - val_accuracy:
 0.3857 - 31s/epoch - 889ms/step

Epoch 2/10

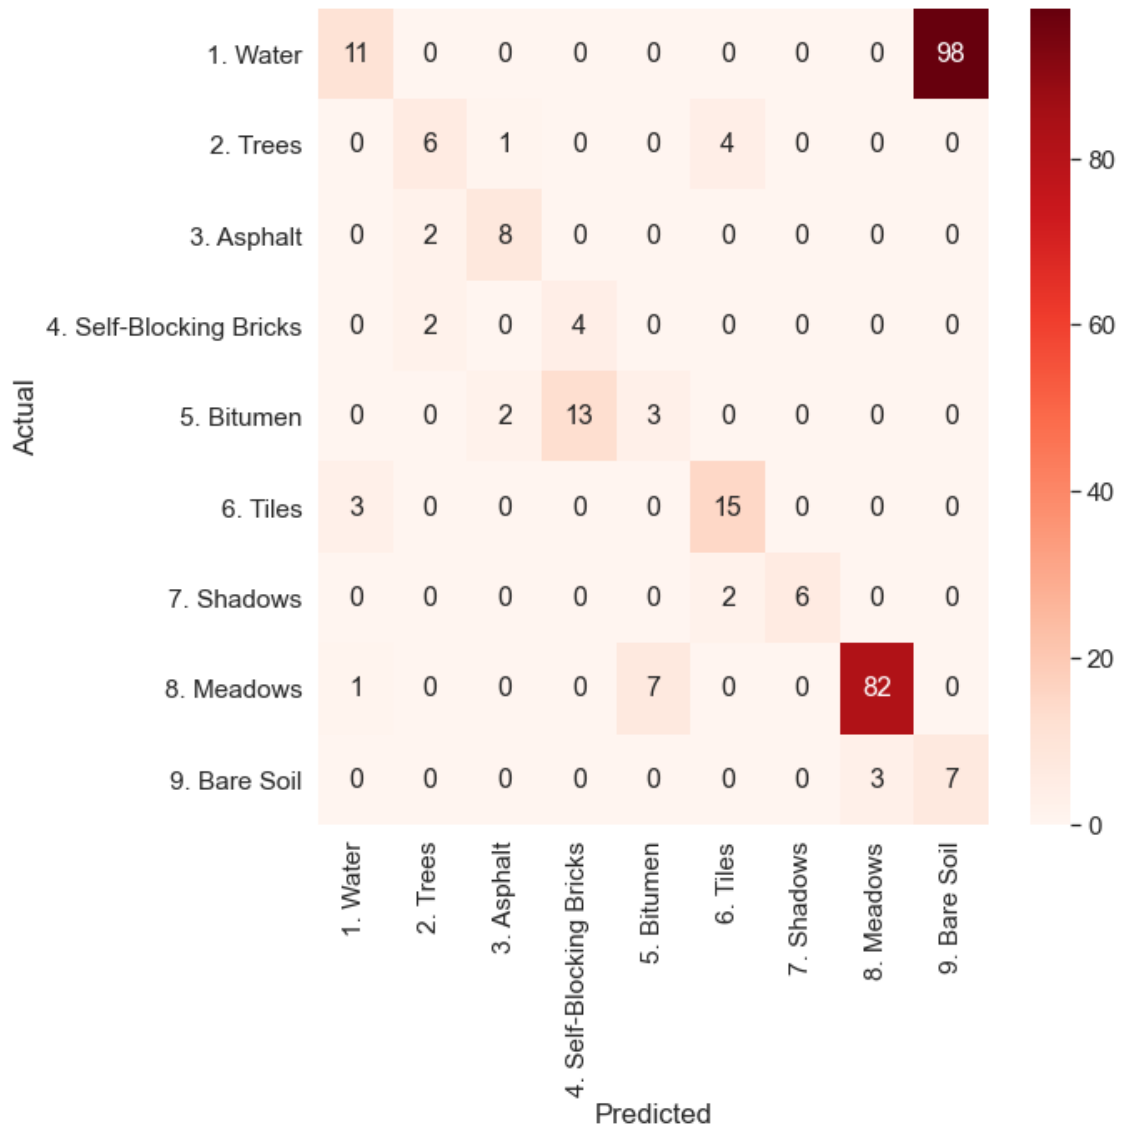
35/35 - 26s - loss: 1.5552 - accuracy: 0.4554 - val_loss: 1.6035 - val_accuracy:
 0.3857 - 26s/epoch - 741ms/step


Epoch 3/10

35/35 - 27s - loss: 1.5538 - accuracy: 0.4554 - val_loss: 1.5756 - val_accuracy:
 0.3857 - 27s/epoch - 777ms/step

Epoch 4/10

35/35 - 27s - loss: 1.5547 - accuracy: 0.4554 - val_loss: 1.5830 - val_accuracy:
0.3857 - 27s/epoch - 769ms/step
Epoch 5/10
35/35 - 25s - loss: 1.5551 - accuracy: 0.4554 - val_loss: 1.5815 - val_accuracy:
0.3857 - 25s/epoch - 706ms/step
Epoch 6/10
35/35 - 26s - loss: 1.5550 - accuracy: 0.4554 - val_loss: 1.5834 - val_accuracy:
0.3857 - 26s/epoch - 754ms/step
Epoch 7/10
35/35 - 28s - loss: 1.5509 - accuracy: 0.4554 - val_loss: 1.5718 - val_accuracy:
0.3857 - 28s/epoch - 795ms/step
Epoch 8/10
35/35 - 27s - loss: 1.5578 - accuracy: 0.4563 - val_loss: 1.5763 - val_accuracy:
0.3857 - 27s/epoch - 778ms/step
Epoch 9/10
35/35 - 27s - loss: 1.5544 - accuracy: 0.4554 - val_loss: 1.5793 - val_accuracy:
0.3857 - 27s/epoch - 763ms/step
Epoch 10/10
35/35 - 34s - loss: 1.5572 - accuracy: 0.4554 - val_loss: 1.5776 - val_accuracy:
0.3857 - 34s/epoch - 972ms/step
<Figure size 432x288 with 0 Axes>





Score for fold 5: loss of 1.5776281356811523; accuracy of 38.57142925262451%
 9/9 [=====] - 1s 43ms/step

```
[[108  0  0  0  0  0  0  0  0]
 [ 18  0  0  0  0  0  0  0  0]
 [  3  0  0  0  0  0  0  0  0]
 [  6  0  0  0  0  0  0  0  0]
 [  8  0  0  0  0  0  0  0  0]
 [ 17  0  0  0  0  0  0  0  0]
 [ 19  0  0  0  0  0  0  0  0]
 [ 99  0  0  0  0  0  0  0  0]
 [  2  0  0  0  0  0  0  0  0]]
```

 Score per fold

 > Fold 1 - Loss: 1.7004808187484741 - Accuracy: 89.64285850524902%

> Fold 2 - Loss: 0.5309717059135437 - Accuracy: 80.0000011920929%

> Fold 3 - Loss: 1.034110188484192 - Accuracy: 62.85714507102966%

> Fold 4 - Loss: 6.309820652008057 - Accuracy: 50.71428418159485%

> Fold 5 - Loss: 1.5776281356811523 - Accuracy: 38.57142925262451%

Average scores for all folds:

> Accuracy: 64.35714364051819 (+- 18.640531349818097)

> Loss: 2.230602300167084

Predicted Overall 1. Water 2. Trees 3. Asphalt \

Actual Overall

1. Water	520	0	0
2. Trees	35	22	3
3. Asphalt	9	10	13
4. Self-Blocking Bricks	7	7	5
5. Bitumen	11	3	6
6. Tiles	59	1	0
7. Shadows	29	6	1
8. Meadows	147	1	0
9. Bare Soil	3	0	0

Predicted Overall 4. Self-Blocking Bricks 5. Bitumen 6. Tiles \

Actual Overall

1. Water	0	0	0
2. Trees	0	5	4
3. Asphalt	0	5	1
4. Self-Blocking Bricks	4	3	0
5. Bitumen	13	19	0
6. Tiles	0	1	30
7. Shadows	0	3	2
8. Meadows	0	9	0
9. Bare Soil	0	0	0

Predicted Overall 7. Shadows 8. Meadows 9. Bare Soil

Actual Overall

1. Water	0	0	98
2. Trees	1	0	0
3. Asphalt	1	1	0
4. Self-Blocking Bricks	1	0	0
5. Bitumen	1	3	0
6. Tiles	0	0	0
7. Shadows	14	0	0
8. Meadows	1	262	0
9. Bare Soil	0	3	17

<Figure size 432x288 with 0 Axes>

