

Full_4 X_Xception_centra

June 13, 2023

1 Date: 9 2022

2 Method: Cross_Inception

3 Data: Pavia

4 Results v.05

```
[ ]: # Libraries
import pandas as pd
import numpy as np
import seaborn as sn
from sklearn.decomposition import PCA
```

```
[ ]: # Read dataset Pavia
from scipy.io import loadmat

def read_HSI():
    X = loadmat('Pavia.mat')['pavia']
    y = loadmat('Pavia_gt.mat')['pavia_gt']
    print(f"X shape: {X.shape}\ny shape: {y.shape}")
    return X, y

X, y = read_HSI()
```

X shape: (1096, 715, 102)

y shape: (1096, 715)

```
[ ]: # PCA
def applyPCA(X, numComponents): # numComponents=64
    newX = np.reshape(X, (-1, X.shape[2]))
    print(newX.shape)
    pca = PCA(n_components=numComponents, whiten=True)
    newX = pca.fit_transform(newX)
    newX = np.reshape(newX, (X.shape[0], X.shape[1], numComponents))
    return newX, pca, pca.explained_variance_ratio_
```

```
[ ]: # channel_wise_shift
def channel_wise_shift(X,numComponents):
    X_copy = np.zeros((X.shape[0] , X.shape[1], X.shape[2]))
    half = int(numComponents/2)
    for i in range(0,half-1):
        X_copy[:, :, i] = X[:, :, (half-i)*2-1]
    for i in range(half,numComponents):
        X_copy[:, :, i] = X[:, :, (i-half)*2]
    X = X_copy
    return X

[ ]: # Split the hyperspectral image into patches of size windowSize-by-windowSize
    ↳pixels
def Patches_Creating(X, y, windowSize, removeZeroLabels = True): #
    ↳windowSize=15, 25
    margin = int((windowSize - 1) / 2)
    zeroPaddedX = padWithZeros(X, margin=margin)
    # split patches
    patchesData = np.zeros((X.shape[0] * X.shape[1], windowSize, windowSize, X.
    ↳shape[2]),dtype="float16")
    patchesLabels = np.zeros((X.shape[0] * X.shape[1]),dtype="float16")
    patchIndex = 0
    for r in range(margin, zeroPaddedX.shape[0] - margin):
        for c in range(margin, zeroPaddedX.shape[1] - margin):
            patch = zeroPaddedX[r - margin:r + margin + 1, c - margin:c +
    ↳margin + 1]
            patchesData[patchIndex, :, :, :] = patch
            patchesLabels[patchIndex] = y[r-margin, c-margin]
            patchIndex = patchIndex + 1
    if removeZeroLabels:
        patchesData = patchesData[patchesLabels>0,:,:,:]
        patchesLabels = patchesLabels[patchesLabels>0]
        patchesLabels -= 1
    return patchesData, patchesLabels
# padding With Zeros
def padWithZeros(X, margin=2):
    newX = np.zeros((X.shape[0] + 2 * margin, X.shape[1] + 2* margin, X.
    ↳shape[2]),dtype="float16")
    x_offset = margin
    y_offset = margin
    newX[x_offset:X.shape[0] + x_offset, y_offset:X.shape[1] + y_offset, :] = X
    return newX

[ ]: # Split Data
from sklearn.model_selection import train_test_split

def splitTrainTestSet(X, y, testRatio, randomState=345):
```

```

X_train, X_test, y_train, y_test = train_test_split(X, y,
↳test_size=testRatio, random_state=randomState,stratify=y)
return X_train, X_test, y_train, y_test

```

```

[ ]: test_ratio = 0.5

# Load and reshape data for training
X0, y0 = read_HSI()
#X=X0
#y=y0

windowSize=15 # accuracy of
# Score for fold 1: loss of 0.34631192684173584; accuracy of 89.49999809265137%

# to test: 7, 9, 13, 15,

width = windowSize
height = windowSize
img_width, img_height, img_num_channels = windowSize, windowSize, 3

input_image_size=windowSize
INPUT_IMG_SIZE=windowSize

dimReduction=3

InputShape=(windowSize, windowSize, dimReduction)

#X, y = loadData(dataset) channel_wise_shift
X1,pca,ratio = applyPCA(X0,numComponents=dimReduction)
X2_shifted = channel_wise_shift(X1,dimReduction) # channel-wise shift
#X2=X1

#print(f"X0 shape: {X0.shape}\ny0 shape: {y0.shape}")
#print(f"X1 shape: {X1.shape}\nX2 shape: {X2.shape}")

X3, y3 = Patches_Creating(X2_shifted, y0, windowSize=windowSize)
Xtrain, Xtest, ytrain, ytest = splitTrainTestSet(X3, y3, test_ratio)

```

```

X shape: (1096, 715, 102)
y shape: (1096, 715)
(783640, 102)

```

```

[ ]: # Compile the model
#incept_model.compile(optimizer='rmsprop', loss='categorical_crossentropy',
↳metrics=['accuracy'])

```

```
[ ]: print()

import warnings
warnings.filterwarnings("ignore")

# load libraries
from keras.initializers import VarianceScaling
from keras.regularizers import l2
from keras.models import Sequential
from keras.layers import Dense
from sklearn import datasets
from sklearn.model_selection import StratifiedKFold
import numpy as np
```

```
[ ]: # 9 classes names

names = ['1. Water', '2. Trees', '3. Asphalt', '4. Self-Blocking Bricks',
         '5. Bitumen', '6. Tiles', '7. Shadows',
         '8. Meadows', '9. Bare Soil']
```

```
[ ]: from tensorflow.keras.applications import EfficientNetB0
from keras.applications import densenet, inception_v3, mobilenet, resnet,
    ↳ vgg16, vgg19, xception
from tensorflow.keras import layers
from keras.layers import Dense, GlobalAveragePooling2D, Dropout, Flatten
import tensorflow as tf

#####
#model = EfficientNetB0(weights='imagenet')

def build_model(num_classes):
    inputs = layers.Input(shape=(windowSize, windowSize, 3))
    #x = img_augmentation(inputs)
    model = xception.Xception(weights='imagenet', include_top=False,
    ↳ input_tensor=inputs)

    #model1 = resnet.ResNet50(weights='imagenet')

    # Freeze the pretrained weights
    model.trainable = False

    # Rebuild top
    x = layers.GlobalAveragePooling2D(name="avg_pool")(model.output)
```

```

x = layers.BatchNormalization()(x)

x = model.output

x = GlobalAveragePooling2D()(x)
# let's add a fully-connected layer
x = Dense(256, activation='relu')(x)
x = Dropout(0.25)(x)

top_dropout_rate = 0.2
#x = layers.Dropout(top_dropout_rate, name="top_dropout")(x)
outputs = layers.Dense(9, activation="softmax", name="pred")(x)

# Compile
model = tf.keras.Model(inputs, outputs, name="EfficientNet")
optimizer = tf.keras.optimizers.Adam(learning_rate=1e-3)
model.compile(
    optimizer=optimizer, loss="categorical_crossentropy",
    metrics=["accuracy"])
return model

```

```

[ ]: '\'\n#model = EfficientNetB0(weights=\'imagenet\')\n\nndef
build_model(num_classes):\n    inputs = layers.Input(shape=(windowSize,
windowSize, 3))\n    #x = img_augmentation(inputs)\n    model =
ception.Xception(weights=\'imagenet\', include_top=False,
input_tensor=inputs)\n\n    #model1 =
resnet.ResNet50(weights=\'imagenet\')\n\n\n    # Freeze the pretrained weights\n
model.trainable = False\n\n    # Rebuild top\n    x =
layers.GlobalAveragePooling2D(name="avg_pool")(model.output)\n    x =
layers.BatchNormalization()(x)\n\n    x = model.output\n\n\n    x =
GlobalAveragePooling2D()(x)\n    # let's add a fully-connected layer\n    x =
Dense(256, activation=\'relu\')(x)\n    x = Dropout(0.25)(x)\n\n\n
top_dropout_rate = 0.2\n    #x = layers.Dropout(top_dropout_rate,
name="top_dropout")(x)\n    outputs = layers.Dense(9, activation="softmax",
name="pred")(x)\n\n    # Compile\n    model = tf.keras.Model(inputs, outputs,
name="EfficientNet")\n    optimizer =
tf.keras.optimizers.Adam(learning_rate=1e-3)\n    model.compile(\n
optimizer=optimizer, loss="categorical_crossentropy", metrics=["accuracy"]\n
)\n    return model\n'

```

```

[ ]: from tensorflow.keras.applications import EfficientNetB0

```

```

def build_model(num_classes):
    inputs = layers.Input(shape=(windowSize, windowSize, 3))
    #x = img_augmentation(inputs)
    #model = EfficientNetB0(include_top=False, input_tensor=inputs,
    ↪weights="imagenet")
    model = xception.Xception(weights='imagenet', include_top=False,
    ↪input_tensor=inputs)

    # Freeze the pretrained weights
    #model.trainable = False

    # Rebuild top
    x = layers.GlobalAveragePooling2D(name="avg_pool")(model.output)
    x = layers.BatchNormalization()(x)

    top_dropout_rate = 0.2
    x = layers.Dropout(top_dropout_rate, name="top_dropout")(x)
    outputs = layers.Dense(9, activation="softmax", name="pred")(x)

    # Compile
    model = tf.keras.Model(inputs, outputs, name="EfficientNet")
    optimizer = tf.keras.optimizers.Adam(learning_rate=1e-3)
    model.compile(
        optimizer=optimizer, loss="categorical_crossentropy",
    ↪metrics=["accuracy"])
    return model

```

```
[ ]: model = build_model(num_classes=9)
```

```

[ ]: def unfreeze_model(model):
    # We unfreeze the top 20 layers while leaving BatchNorm layers frozen
    for layer in model.layers[-20:]:
        if not isinstance(layer, layers.BatchNormalization):
            layer.trainable = True

    optimizer = tf.keras.optimizers.Adam(learning_rate=1e-4)
    model.compile(
        optimizer=optimizer, loss="categorical_crossentropy",
    ↪metrics=["accuracy"])

```

```
[ ]: import matplotlib.pyplot as plt
```

```
def plot_hist(hist):
```

```

plt.plot(hist.history["accuracy"])
plt.plot(hist.history["val_accuracy"])
plt.title("model accuracy")
plt.ylabel("accuracy")
plt.xlabel("epoch")
plt.legend(["train", "validation"], loc="upper left")
plt.show()

```

```

[ ]: from tensorflow.keras.losses import sparse_categorical_crossentropy
from tensorflow.keras.optimizers import Adam
from sklearn.model_selection import KFold
from tensorflow.keras import layers

import numpy as np
from sklearn.metrics import confusion_matrix, accuracy_score, \
    ↪classification_report, cohen_kappa_score
import matplotlib.pyplot as plt
from keras.applications.inception_resnet_v2 import InceptionResNetV2, \
    ↪preprocess_input
from keras.layers import Dense, GlobalAveragePooling2D, Dropout, Flatten
from keras.models import Model

import tensorflow as tf

# configuration
confmat = 0
batch_size = 50
loss_function = sparse_categorical_crossentropy
no_classes = 9
no_epochs = 10
optimizer = Adam()
verbosity = 1
num_folds = 5

NN=len(Xtrain)
#NN=2000
NN=5000

input_train=Xtrain[0:NN]
target_train=ytrain[0:NN]

input_test=Xtest[0:NN]
target_test=ytest[0:NN]

# Determine shape of the data
input_shape = (img_width, img_height, img_num_channels)

```

```

# Parse numbers as floats
_train = _train.astype('float32')
_test = _test.astype('float32')

# Normalize data
_train = _train / 255
_test = _test / 255

# Define per-fold score containers
acc_per_fold = []
loss_per_fold = []

Y_pred=[]
y_pred=[]
# Merge inputs and targets
inputs = np.concatenate((_train, _test), axis=0)
targets = np.concatenate((target_train, target_test), axis=0)

# Define the K-fold Cross Validator
kfold = KFold(n_splits=num_folds, shuffle=True)

# K-fold Cross Validation model evaluation
fold_no = 1
for train, test in kfold.split(inputs, targets):

    # model architecture

    # Compile the model
    #model.compile(optimizer='rmsprop', loss='categorical_crossentropy',
    →metrics=['accuracy'])

    # Compile the model
    # model.compile(optimizer='rmsprop', loss='categorical_crossentropy',
    →metrics=['accuracy'])

    model = build_model(num_classes=9)
    #model.compile(loss=loss_function, optimizer='rmsprop',metrics=['accuracy'])

    #model.summary()

    #unfreeze_model(model)
    model.compile(loss=loss_function, optimizer='rmsprop',metrics=['accuracy'])

```



```

# Generate a print
↳
↳ print('-----')
print(f'Training for fold {fold_no} ...')

# Fit data to model
#model.summary()

history = model.fit(inputs[train], targets[train],
                    validation_data = (inputs[test], targets[test]),
                    epochs=no_epochs, verbose=2 )
plt.figure()
plot_hist(history)
# hist = model.fit(inputs[train], targets[train],
#                  steps_per_epoch=(29943/batch_size),
#                  epochs=5,
#                  validation_data=(inputs[test], targets[test]),
#                  validation_steps=(8000/batch_size),
#                  initial_epoch=20,
#                  verbose=1 )
plt.figure()

# Generate generalization metrics
scores = model.evaluate(inputs[test], targets[test], verbose=0)
print(f'Score for fold {fold_no}: {model.metrics_names[0]} of {scores[0]};
↳ {model.metrics_names[1]} of {scores[1]*100}%')
acc_per_fold.append(scores[1] * 100)
loss_per_fold.append(scores[0])

# confusion_matrix
Y_pred = model.predict(inputs[test])
y_pred = np.argmax(Y_pred, axis=1)
#target_test=targets[test]

confusion = confusion_matrix(targets[test], y_pred)
df_cm = pd.DataFrame(confusion, columns=np.unique(names), index = np.
↳ unique(names))
df_cm.index.name = 'Actual'
df_cm.columns.name = 'Predicted'
plt.figure(figsize = (9,9))
sn.set(font_scale=1.4)#for label size
sn.heatmap(df_cm, cmap="Reds", annot=True, annot_kws={"size": 16}, fmt='d')
plt.savefig('cmap.png', dpi=300)
print(confusion_matrix(targets[test], y_pred))

```

```

confmat      = confmat + confusion;

# Increase fold number
fold_no = fold_no + 1

# == average scores ==
print('-----')
print('Score per fold')
for i in range(0, len(acc_per_fold)):
    ↪print('-----')
    ↪print(f'> Fold {i+1} - Loss: {loss_per_fold[i]} - Accuracy:↪
    ↪{acc_per_fold[i]}%')
print('-----')
print('Average scores for all folds:')
print(f'> Accuracy: {np.mean(acc_per_fold)} (+- {np.std(acc_per_fold)})')
print(f'> Loss: {np.mean(loss_per_fold)}')
print('-----')

Overall_Conf = pd.DataFrame(confmat, columns=np.unique(names), index = np.
    ↪unique(names))
Overall_Conf.index.name = 'Actual Overall'
Overall_Conf.columns.name = 'Predicted Overall'
plt.figure(figsize = (10,8))
sn.set(font_scale=1.4)#for label size
sn.heatmap(Overall_Conf, cmap="Reds", annot=True,annot_kws={"size": 16},↪
    ↪fmt='d')
plt.savefig('cmap.png', dpi=300)
print(Overall_Conf)

# Notes for next trial

# window size=25 __> will work
# window size=5 --> Only Bayesian will work
# Need to test (7, 9, 11, 13, 15) window sizes
# When the accuracy is decreasing, it's not right.
# When need to get acc over 0.7

```

Training for fold 1 ...


Epoch 1/10

250/250 - 207s - loss: 0.8361 - accuracy: 0.7336 - val_loss: 0.9933 -
val_accuracy: 0.7275 - 207s/epoch - 830ms/step

Epoch 2/10

250/250 - 201s - loss: 0.3387 - accuracy: 0.9075 - val_loss: 0.1749 -

val_accuracy: 0.9680 - 201s/epoch - 806ms/step
Epoch 3/10
250/250 - 241s - loss: 0.2152 - accuracy: 0.9426 - val_loss: 6.2874 -
val_accuracy: 0.6280 - 241s/epoch - 962ms/step
Epoch 4/10
250/250 - 252s - loss: 0.1856 - accuracy: 0.9539 - val_loss: 1.7030 -
val_accuracy: 0.9315 - 252s/epoch - 1s/step
Epoch 5/10
250/250 - 256s - loss: 0.1201 - accuracy: 0.9647 - val_loss: 0.3575 -
val_accuracy: 0.9665 - 256s/epoch - 1s/step
Epoch 6/10
250/250 - 211s - loss: 0.0937 - accuracy: 0.9711 - val_loss: 0.1008 -
val_accuracy: 0.9710 - 211s/epoch - 844ms/step
Epoch 7/10
250/250 - 211s - loss: 0.0867 - accuracy: 0.9772 - val_loss: 0.0528 -
val_accuracy: 0.9860 - 211s/epoch - 845ms/step
Epoch 8/10
250/250 - 215s - loss: 0.0828 - accuracy: 0.9801 - val_loss: 0.1079 -
val_accuracy: 0.9725 - 215s/epoch - 860ms/step
Epoch 9/10
250/250 - 211s - loss: 0.0620 - accuracy: 0.9837 - val_loss: 0.0815 -
val_accuracy: 0.9775 - 211s/epoch - 843ms/step
Epoch 10/10
250/250 - 186s - loss: 0.0590 - accuracy: 0.9844 - val_loss: 0.9210 -
val_accuracy: 0.8710 - 186s/epoch - 743ms/step



Score for fold 1: loss of 0.9210289716720581; accuracy of 87.09999918937683%
63/63 [=====] - 2s 28ms/step

```
[[667  0  0 29  0  0  0  0 207]
 [ 0 71  0  0  0  0  0  0  0]
 [ 0  4 37  2  0  1  2  0  0]
 [ 0  0  0 35  0  0  0  0  0]
 [ 1  0  0  2 103  0  0  0  0]
 [ 0  2  0  0  0 120  1  0  0]
 [ 0  3  0  0  0  0 83  0  0]
 [ 0  0  0  3  0  0  0 585  0]
 [ 0  0  0  0  0  1  0  0 41]]
```

Training for fold 2 ...

Epoch 1/10

250/250 - 188s - loss: 1.4509 - accuracy: 0.5096 - val_loss: 1.3032 -
val_accuracy: 0.6645 - 188s/epoch - 751ms/step

Epoch 2/10

250/250 - 172s - loss: 0.3862 - accuracy: 0.8857 - val_loss: 13.8235 -
val_accuracy: 0.5430 - 172s/epoch - 687ms/step

Epoch 3/10

250/250 - 176s - loss: 0.2133 - accuracy: 0.9371 - val_loss: 0.3055 -
val_accuracy: 0.9285 - 176s/epoch - 702ms/step

Epoch 4/10

250/250 - 161s - loss: 0.1449 - accuracy: 0.9559 - val_loss: 2.6323 -
val_accuracy: 0.9070 - 161s/epoch - 646ms/step

Epoch 5/10

250/250 - 165s - loss: 0.1310 - accuracy: 0.9609 - val_loss: 2.0419 -
val_accuracy: 0.9620 - 165s/epoch - 660ms/step

Epoch 6/10

250/250 - 167s - loss: 0.1523 - accuracy: 0.9647 - val_loss: 0.1607 -
val_accuracy: 0.9780 - 167s/epoch - 667ms/step

Epoch 7/10

250/250 - 180s - loss: 0.0952 - accuracy: 0.9736 - val_loss: 0.0548 -
val_accuracy: 0.9840 - 180s/epoch - 719ms/step

Epoch 8/10

250/250 - 209s - loss: 0.0845 - accuracy: 0.9768 - val_loss: 0.0547 -
val_accuracy: 0.9875 - 209s/epoch - 835ms/step


Epoch 9/10


250/250 - 212s - loss: 0.0653 - accuracy: 0.9836 - val_loss: 0.1220 -
val_accuracy: 0.9695 - 212s/epoch - 848ms/step

Epoch 10/10

250/250 - 203s - loss: 0.0674 - accuracy: 0.9830 - val_loss: 0.0432 -
val_accuracy: 0.9860 - 203s/epoch - 811ms/step

<Figure size 432x288 with 0 Axes>





Score for fold 2: loss of 0.043157752603292465; accuracy of 98.60000014305115%
 63/63 [=====] - 3s 35ms/step

```
[[975  0  0  0  0  0  0  0  0]
 [  0 96  7  4  0  0  0  0  0]
 [  0  1 27  1  9  0  0  0  0]
 [  0  0  0 34  0  0  0  0  0]
 [  0  0  0  0 83  0  0  0  0]
 [  0  1  1  0  1 110  1  0  0]
 [  0  0  0  0  0  0 80  0  0]
 [  0  0  0  0  0  0  0 534  0]
 [  0  0  0  0  0  0  1  1 33]]
```

 Training for fold 3 ...

Epoch 1/10

250/250 - 225s - loss: 0.7617 - accuracy: 0.7688 - val_loss: 0.7705 -
 val_accuracy: 0.8000 - 225s/epoch - 900ms/step

Epoch 2/10

250/250 - 214s - loss: 0.2617 - accuracy: 0.9216 - val_loss: 0.2535 -
 val_accuracy: 0.9485 - 214s/epoch - 857ms/step


Epoch 3/10


250/250 - 207s - loss: 0.1870 - accuracy: 0.9467 - val_loss: 0.0707 -
 val_accuracy: 0.9765 - 207s/epoch - 828ms/step

Epoch 4/10

250/250 - 186s - loss: 0.1347 - accuracy: 0.9639 - val_loss: 0.0912 -
val_accuracy: 0.9735 - 186s/epoch - 744ms/step
Epoch 5/10
250/250 - 181s - loss: 0.0951 - accuracy: 0.9720 - val_loss: 0.0523 -
val_accuracy: 0.9850 - 181s/epoch - 722ms/step
Epoch 6/10
250/250 - 184s - loss: 0.1060 - accuracy: 0.9730 - val_loss: 6.1422 -
val_accuracy: 0.6715 - 184s/epoch - 738ms/step
Epoch 7/10
250/250 - 195s - loss: 0.0784 - accuracy: 0.9793 - val_loss: 0.0569 -
val_accuracy: 0.9770 - 195s/epoch - 781ms/step
Epoch 8/10
250/250 - 196s - loss: 0.0687 - accuracy: 0.9830 - val_loss: 0.4723 -
val_accuracy: 0.9635 - 196s/epoch - 786ms/step
Epoch 9/10
250/250 - 181s - loss: 0.0503 - accuracy: 0.9862 - val_loss: 3.6070 -
val_accuracy: 0.8265 - 181s/epoch - 725ms/step
Epoch 10/10
250/250 - 177s - loss: 0.0433 - accuracy: 0.9875 - val_loss: 0.0637 -
val_accuracy: 0.9825 - 177s/epoch - 708ms/step

<Figure size 432x288 with 0 Axes>





Score for fold 3: loss of 0.06372090429067612; accuracy of 98.25000166893005%
 63/63 [=====] - 2s 30ms/step

```
[[884  0  0  0  0  0  0  0  0]
 [  0 96  4  0  0  0  0  0  0]
 [  0  1 43  0  3  1  0  0  0]
 [  0  4  2 34  2  0  0  0  0]
 [  0  2  1  0 92  0  0  0  0]
 [  0  0  1  0  0 141  0  1  0]
 [  0  2  2  0  0  3  73  0  0]
 [  1  1  0  0  0  0  0 553  2]
 [  2  0  0  0  0  0  0  0 49]]
```

 Training for fold 4 ...

Epoch 1/10

250/250 - 183s - loss: 0.7109 - accuracy: 0.7980 - val_loss: 0.4502 -
 val_accuracy: 0.8890 - 183s/epoch - 733ms/step

Epoch 2/10

250/250 - 185s - loss: 0.3008 - accuracy: 0.9101 - val_loss: 13.0653 -
 val_accuracy: 0.9285 - 185s/epoch - 741ms/step

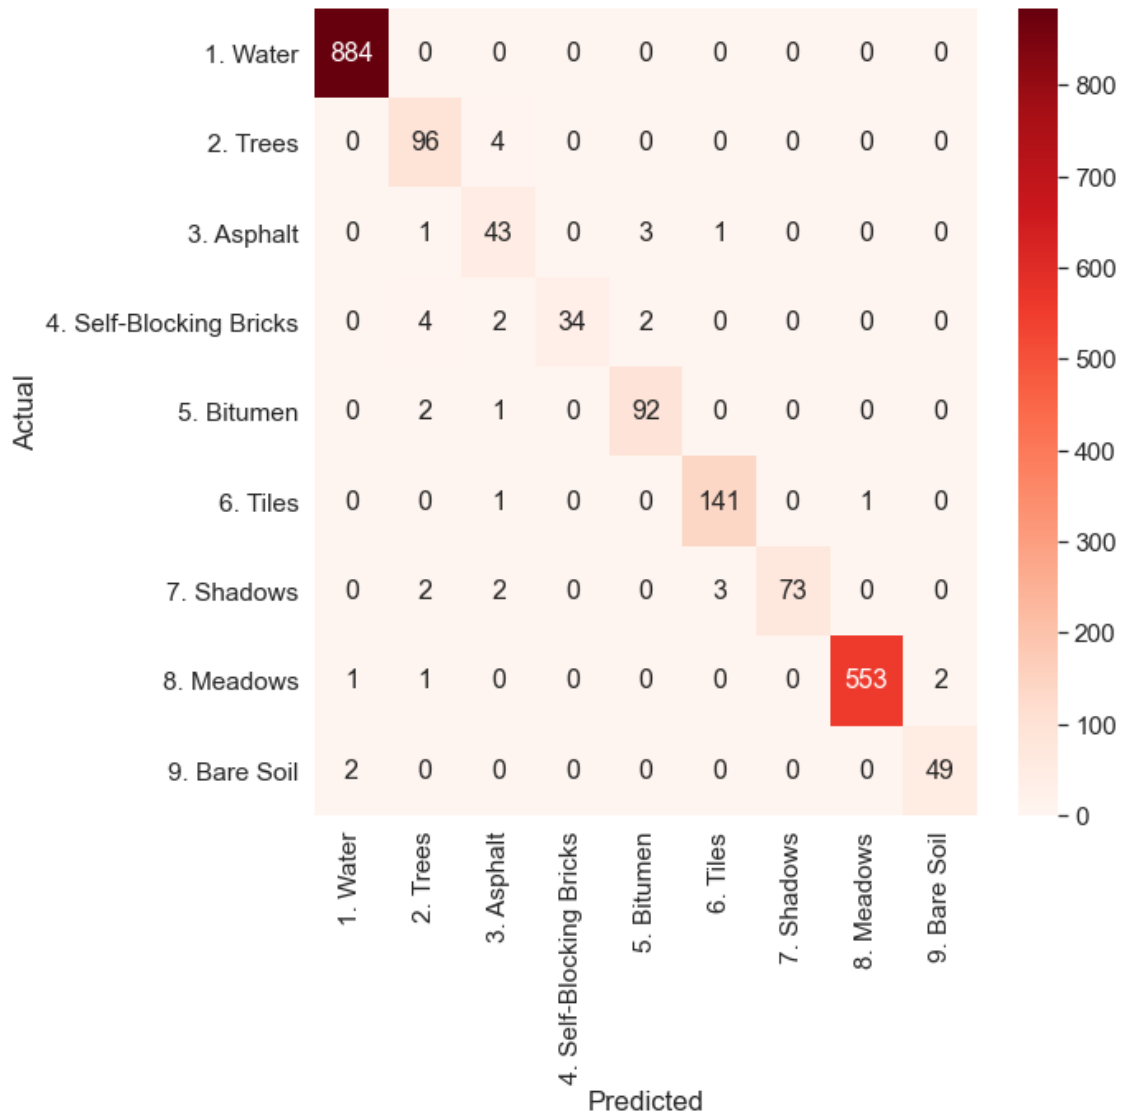
Epoch 3/10


250/250 - 171s - loss: 0.1945 - accuracy: 0.9442 - val_loss: 0.7509 -
 val_accuracy: 0.9365 - 171s/epoch - 683ms/step

Epoch 4/10

250/250 - 214s - loss: 0.1762 - accuracy: 0.9511 - val_loss: 0.1980 -
val_accuracy: 0.9370 - 214s/epoch - 856ms/step
Epoch 5/10
250/250 - 205s - loss: 0.1123 - accuracy: 0.9660 - val_loss: 0.0768 -
val_accuracy: 0.9735 - 205s/epoch - 821ms/step
Epoch 6/10
250/250 - 219s - loss: 0.1042 - accuracy: 0.9729 - val_loss: 1.0592 -
val_accuracy: 0.9100 - 219s/epoch - 874ms/step
Epoch 7/10
250/250 - 221s - loss: 0.0930 - accuracy: 0.9765 - val_loss: 0.0700 -
val_accuracy: 0.9810 - 221s/epoch - 884ms/step
Epoch 8/10
250/250 - 194s - loss: 0.0598 - accuracy: 0.9844 - val_loss: 0.0287 -
val_accuracy: 0.9905 - 194s/epoch - 777ms/step
Epoch 9/10
250/250 - 198s - loss: 0.0713 - accuracy: 0.9827 - val_loss: 0.0545 -
val_accuracy: 0.9850 - 198s/epoch - 794ms/step
Epoch 10/10
250/250 - 195s - loss: 0.0528 - accuracy: 0.9862 - val_loss: 0.0690 -
val_accuracy: 0.9785 - 195s/epoch - 782ms/step

<Figure size 432x288 with 0 Axes>





Score for fold 4: loss of 0.0689583271741867; accuracy of 97.85000085830688%
 63/63 [=====] - 2s 32ms/step

```
[[899  0  1  0  0  0  0  0  0]
 [  0 77  6  7  0  0  1  0  0]
 [  0  1 30  0  3  0  0  0  0]
 [  0  0  0 39  0  0  0  0  0]
 [  0  0  3  0 77  0  0  0  0]
 [  0  0 14  0  0 127  1  0  0]
 [  0  0  0  1  0  0  99  0  0]
 [  0  0  0  1  0  0  0 575  0]
 [  4  0  0  0  0  0  0  0 34]]
```

 Training for fold 5 ...

Epoch 1/10

250/250 - 201s - loss: 0.8894 - accuracy: 0.7170 - val_loss: 1.2928 -
 val_accuracy: 0.7270 - 201s/epoch - 805ms/step

Epoch 2/10

250/250 - 195s - loss: 0.2875 - accuracy: 0.9126 - val_loss: 4.7491 -
 val_accuracy: 0.5755 - 195s/epoch - 780ms/step


Epoch 3/10


250/250 - 196s - loss: 0.2189 - accuracy: 0.9352 - val_loss: 12.2001 -
 val_accuracy: 0.5310 - 196s/epoch - 782ms/step

Epoch 4/10

250/250 - 195s - loss: 0.1664 - accuracy: 0.9538 - val_loss: 0.0718 -
val_accuracy: 0.9770 - 195s/epoch - 782ms/step
Epoch 5/10
250/250 - 196s - loss: 0.1092 - accuracy: 0.9681 - val_loss: 0.0463 -
val_accuracy: 0.9850 - 196s/epoch - 784ms/step
Epoch 6/10
250/250 - 195s - loss: 0.1175 - accuracy: 0.9689 - val_loss: 0.2958 -
val_accuracy: 0.9770 - 195s/epoch - 782ms/step
Epoch 7/10
250/250 - 196s - loss: 0.0760 - accuracy: 0.9768 - val_loss: 0.0466 -
val_accuracy: 0.9885 - 196s/epoch - 783ms/step
Epoch 8/10
250/250 - 195s - loss: 0.0865 - accuracy: 0.9784 - val_loss: 0.1726 -
val_accuracy: 0.9745 - 195s/epoch - 782ms/step
Epoch 9/10
250/250 - 195s - loss: 0.0871 - accuracy: 0.9781 - val_loss: 0.0523 -
val_accuracy: 0.9890 - 195s/epoch - 781ms/step
Epoch 10/10
250/250 - 194s - loss: 0.0651 - accuracy: 0.9836 - val_loss: 0.0530 -
val_accuracy: 0.9870 - 194s/epoch - 778ms/step

<Figure size 432x288 with 0 Axes>





Score for fold 5: loss of 0.0529952198266983; accuracy of 98.69999885559082%
 63/63 [=====] - 3s 34ms/step

```
[[919  0  0  0  0  0  0  0  0]
 [  0 99  1  0  0  0  1  0  0]
 [  0  8 35  0  0  0  0  1  0]
 [  0  0  0 41  0  0  0  0  0]
 [  0  0  1  7 75  0  0  0  0]
 [  2  1  0  0  0 101  1  0  0]
 [  0  0  0  0  0  0  96  0  0]
 [  0  0  0  0  0  0  0 575  1]
 [  2  0  0  0  0  0  0  0 33]]
```

 Score per fold

 > Fold 1 - Loss: 0.9210289716720581 - Accuracy: 87.09999918937683%

> Fold 2 - Loss: 0.043157752603292465 - Accuracy: 98.60000014305115%

> Fold 3 - Loss: 0.06372090429067612 - Accuracy: 98.25000166893005%

> Fold 4 - Loss: 0.0689583271741867 - Accuracy: 97.85000085830688%

> Fold 5 - Loss: 0.0529952198266983 - Accuracy: 98.69999885559082%

Average scores for all folds:

> Accuracy: 96.10000014305115 (+- 4.509878480314604)


> Loss: 0.22997223511338233

Predicted Overall	1. Water	2. Trees	3. Asphalt \
Actual Overall			
1. Water	4344	0	1
2. Trees	0	439	18
3. Asphalt	0	15	172
4. Self-Blocking Bricks	0	4	2
5. Bitumen	1	2	5
6. Tiles	2	4	16
7. Shadows	0	5	2
8. Meadows	1	1	0
9. Bare Soil	8	0	0

Predicted Overall	4. Self-Blocking Bricks	5. Bitumen	6. Tiles \
Actual Overall			
1. Water	29	0	0
2. Trees	11	0	0
3. Asphalt	3	15	2
4. Self-Blocking Bricks	183	2	0
5. Bitumen	9	430	0
6. Tiles	0	1	599
7. Shadows	1	0	3
8. Meadows	4	0	0
9. Bare Soil	0	0	1

Predicted Overall	7. Shadows	8. Meadows	9. Bare Soil
Actual Overall			
1. Water	0	0	207
2. Trees	2	0	0
3. Asphalt	2	1	0
4. Self-Blocking Bricks	0	0	0
5. Bitumen	0	0	0
6. Tiles	4	1	0
7. Shadows	431	0	0
8. Meadows	0	2822	3
9. Bare Soil	1	1	190

<Figure size 432x288 with 0 Axes>





```
[ ]: print(classification_report(targets[test], y_pred, target_names = names))
```

	precision	recall	f1-score	support
1. Water	1.00	1.00	1.00	919
2. Trees	0.92	0.98	0.95	101
3. Asphalt	0.95	0.80	0.86	44
4. Self-Blocking Bricks	0.85	1.00	0.92	41
5. Bitumen	1.00	0.90	0.95	83
6. Tiles	1.00	0.96	0.98	105
7. Shadows	0.98	1.00	0.99	96
8. Meadows	1.00	1.00	1.00	576
9. Bare Soil	0.97	0.94	0.96	35
accuracy			0.99	2000
macro avg	0.96	0.95	0.96	2000
weighted avg	0.99	0.99	0.99	2000